

Image Analysis, Assignment 1

1 Image sampling

Consider a continuous monochrome image. The intensity (brightness) in a point (x, y) in the image is given by the function

$$f(x, y) = x(1 - y), \quad 0 \leq x \leq 1, 0 \leq y \leq 1.$$

Think about (or plot) what the image looks like.

Sample the image evenly to a discrete image with 7×7 pixels. Let the lower left pixel be a sample from the point $(0, 0)$ in the continuous image and the upper right pixel a sample from $(1, 1)$. Quantify the discrete image with 16 different gray levels from 0 to 15, with 0 corresponding to 0, and 15 corresponding to 1 in the continuous image. What is the result?

For the report: Write out the resulting 7×7 image matrix. Explain how you did your calculations.

2 Histogram equalization

An image (in a continuous representation) has gray level histogram

$$p_r = 6r(1 - r), \quad r \in [0, 1].$$

What gray level transform $s = T(r)$ should be used so that the resulting histogram p_s is uniform, i.e.

$$p_s = 1, \quad s \in [0, 1] ?$$

For the report: Specify the transformation $s = T(r)$ and show how you computed it.

3 Neighbourhood of pixels

Consider the following image

$$f = \begin{pmatrix} 3 & 3 & 3 & 3 & 3 & 2 & 2 & 2 & 2 & 2 & 3 & 3 & 3 & 3 & 3 \\ 3 & 3 & 1 & 2 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 2 & 3 & 3 & 3 \\ 3 & 3 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 2 & 3 & 3 & 3 \\ 3 & 2 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 2 & 3 \\ 3 & 1 & 0 & 2 & 3 & 1 & 0 & 0 & 0 & 3 & 3 & 1 & 0 & 1 & 3 \\ 2 & 0 & 0 & 3 & 3 & 2 & 0 & 0 & 1 & 3 & 3 & 2 & 0 & 0 & 2 \\ 2 & 0 & 0 & 2 & 3 & 1 & 0 & 0 & 0 & 2 & 3 & 1 & 0 & 0 & 2 \\ 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 2 \\ 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 2 \\ 2 & 0 & 1 & 2 & 1 & 1 & 0 & 0 & 0 & 1 & 1 & 2 & 2 & 0 & 2 \\ 3 & 1 & 0 & 2 & 3 & 3 & 3 & 3 & 3 & 3 & 3 & 0 & 1 & 3 & 3 \\ 3 & 2 & 0 & 0 & 2 & 3 & 3 & 3 & 3 & 3 & 2 & 0 & 0 & 2 & 3 \\ 3 & 3 & 2 & 0 & 0 & 2 & 3 & 3 & 3 & 3 & 2 & 0 & 0 & 2 & 3 \\ 3 & 3 & 3 & 2 & 1 & 0 & 0 & 0 & 0 & 1 & 2 & 3 & 3 & 3 & 3 \\ 3 & 3 & 3 & 3 & 3 & 2 & 2 & 2 & 2 & 2 & 3 & 3 & 3 & 3 & 3 \end{pmatrix} .$$

- a) Mark each element that has intensity 0 or 1 with a circle. Threshold the image to get a new image g so that $g(i,j) = 0$ if $f(i,j) \leq 1$ and $g(i,j) = 1$ if $f(i,j) > 1$.
- b) Find the 4-connected components for $g = 1$ (as defined in the lectures).

For the report: Give the resulting g with marked connected components. Do this by hand, you should not use the built-in functions for this task.

4 Segmentation part of OCR

During all of the four assignments you are gradually going to build and test a small system for OCR (optical character recognition). During this first assignment your task is to write a function `im2segment`, and to test this function on a few images using a benchmark script . On the Canvas page for the course there is a zip-file with data. By downloading and unpacking the file you will obtain a folder `ocr_project` and in this folder there is a folder

`datasets`

In the `datasets` folder there is for now only one folder 'short1', which contains a few test examples and ground-truth both for segmentation and for recognition. Later on we will add more folders with additional (and more challenging images).

Each image in the folder contains text (light against dark background) Write a function `im2segment` that takes such an image matrix I as input and returns a segmentation, i.e. a set of images $S = (S_1, \dots, S_n)$ one for each digit in the image. Each such image matrix S_i should be a matrix with ones at the pixels for that digit and zeros for all other pixels. Also make sure that each individual output image has the same size as the original image.

Test your function `im2segment` using the benchmark script/function, that loads each image in a folder and each ground truth segmentation and measures how well it works.

For the report:

- Print your code for your implementation of `im2segment`, and the text results of running the benchmark script.
- Include a figure of one example of an input image and the resulting segmentation images, similar to the ones above.
- Make sure that the code is commented so that it is easy to follow, or write a description of what it does.

Matlab-specific instructions

In the OCR folder there is a folder

`matlab`

Study the script `inl1_stub.m`, that reads one of the images in the folder `short1`.

In order to make things easy for all of us, you should all use the same convention to let the output be a so called cell array in Matlab.

```
S = cell(1,n);
S{1} = bild1;
...
S{n} = bildn;
```

Also make sure that each image (e.g. `bild1`) has the same size as the original image. You should implement the function `S = im2segment(image)` in a file `im2segment.m` in Matlab (there is already a very bad implementation in the folder). Test your function using the benchmark script `inl1_test_and_benchmark`.

Python-specific instructions

In the OCR folder there is a script `assignment1.py`. Study the script that reads an example image and runs the function `im2segment` on it. It also runs a benchmark function `benchmark_assignment1` that tests the segmentation function on the dataset. There exists a bad implementation of `im2segment` already, but you should replace this with your own implementation. In order to make things easy for all of us, you should all use the same convention to let the output be a list of numpy arrays `[S[0], S[1], ..., S[n-1]]` containing the `n` segmentations. Also make sure that each image (e.g. `S[k]`) has the same size as the original image. Test your implementation using the benchmark function `benchmark_assignment1`.

Example result

If all goes well when running your segmentation on the following image (im)



the resulting output array will then contain 5 images, i.e. $S\{1\}$ (Matlab) or $S[0]$ (Python) is



the second segment $S\{2\}$ (Matlab) or $S[1]$ (Python) is



and so on. Again notice that each segmented image is of the same size as the original image.

5 Dimensionality

A **vector space** is a collection of objects and two operations **addition** and **multiplication by scalar**. For finite dimensional vector spaces one may choose a basis of elements e_1, \dots, e_k so that every example u can be written as

$$u = x_1e_1 + \dots + x_ke_k$$

for some set of scalars (x_1, \dots, x_k) . Here k is the dimension of the vector space.

A: The set of gray-scale (monochrome) images with 2×4 pixels is a vector space. It has a finite dimension k , and we can define a basis that consists of k 3×2 images.

B: The set of gray-scale (monochrome) images with 1200×3000 pixels is a vector space.

For the report:

- In A what is the dimension k ?
- Give an example basis for A (by explicitly writing down the basis elements e_1, \dots, e_k).
- In B what is the dimension k ?
- In B describe how the basis elements can be chosen, i.e. give the first couple of basis elements and describe how the following ones are constructed.

6 Scalar products and norm on images

Given three images

$$u = \begin{bmatrix} 4 & -3 \\ 2 & -1 \end{bmatrix},$$

$$v = \frac{1}{2} \begin{bmatrix} 1 & -1 \\ -1 & 1 \end{bmatrix},$$

$$w = \frac{1}{2} \begin{bmatrix} 1 & -1 \\ 1 & -1 \end{bmatrix}.$$

For the report:

- How is the scalar product defined for images?
- How is the norm of an image defined?
- Calculate $\|u\|$, $\|v\|$, $\|w\|$, $u \cdot v$, $u \cdot w$, $v \cdot w$.
- Are the matrices $\{v, w\}$ orthonormal?
- What is the orthogonal projection of u on the subspace spanned by $\{v, w\}$? Is the resulting projection a good approximation, i.e. is the projection close/similar to u ?

(Notice that there are many possible matrix norms! There is a potential pitfall here if you use a computer program, e.g. matlab. You might accidentally be using the wrong norm!)

7 Image compression

A small camera delivers low resolution images with 3×4 pixels. Before transmitting the image to a computer, one would like to compress the images consisting of twelve intensities to four numbers. After studying numerous images and using principal component analysis one has determined that the following four images represent typical images well,

$$\phi_1 = \frac{1}{3} \begin{pmatrix} 1 & 1 & 1 \\ 1 & 0 & 1 \\ -1 & -1 & -1 \\ 0 & -1 & 0 \end{pmatrix}, \quad \phi_2 = \frac{1}{2} \begin{pmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{pmatrix}, \quad \phi_3 = \frac{1}{2} \begin{pmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{pmatrix}, \quad \phi_4 = \frac{1}{3} \begin{pmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 1 \end{pmatrix}.$$

Assume that we have an original image

$$f = \begin{pmatrix} 11 & 5 & -3 \\ 12 & 1 & -1 \\ 1 & -4 & -6 \\ 5 & -3 & -2 \end{pmatrix}.$$

We would like to find an approximation f_a of f using the four basis images.

For the report:

- Show that the four images $\phi_1, \phi_2, \phi_3, \phi_4$ are orthonormal (write out all combinations you test, and their result).
- How do you determine the four numbers (coordinates) x_1, x_2, x_3, x_4 such that the approximate image
$$f_a = x_1\phi_1 + x_2\phi_2 + x_3\phi_3 + x_4\phi_4$$
is as close to f as possible, i.e. such that $|f - f_a|^2 = (f - f_a) \cdot (f - f_a)$ is as small as possible?
- Calculate x_1, x_2, x_3, x_4 for the image f above.
- Determine the approximate f_a .
- Is f_a close to f ? Can we expect better or worse approximations than in task 6? Why? Explain your reasoning!

8 Image bases

In the dataset file for the assignment on Canvas there is a file, `assignment1bases.mat` with two variables `bases` and `stacks`.

The variable `stacks` is a cell array. It contains two stacks of images

- 400 general test images of size 19×19 . These are stored in a variable `stacks{1}`, which is a three-dimensional data structure of size $19 \times 19 \times 400$.
- 400 test images of faces of size 19×19 . These are stored in a variable `stacks{2}`, which is a three-dimensional data structure of size $19 \times 19 \times 400$.

In the previous exercise we saw how to project an image onto a low-dimensional subspace defined by a set of basis images. Some interesting questions for discussion are:

- Is there a difference between different bases?
- What is the best basis?
- How can one calculate a good basis?

We do not expect a full answer to these three questions, but we encourage you to think about them and to discuss them with your fellow students.

The variable `bases` is also a cell array. It contains three sets of bases for three different subspaces of dimension four. The first basis is stored in a variable `bases{1}`, which is a tensor of size $19 \times 19 \times 4$. Thus the four basis images are `bases{1}(:,:,1)`, `bases{1}(:,:,2)`, `bases{1}(:,:,3)`, `bases{1}(:,:,4)`.

Write a function that projects an image u onto a basis (e_1, e_2, e_3, e_4) and returns the projection u_p and error norm r , i.e. the norm of the difference $r = |u - u_p|$.

Then write a script that tests all of the 400 test images in a test set and returns the mean of the error norms. Calculate this mean for each of the two test sets on each of the three bases.

Discussion—How should one calculate the best basis for a stack of images?

Python: You can load the variables `bases` and `stacks` as numpy-arrays using `scipy.io.loadmat`.

```
import scipy
data = scipy.io.loadmat('assignment1bases.mat')
bases = data['bases']
stacks = data['stacks']
```

For the report:

- Include the code of your function for projection, with explanation of how it works.
- plot a few images in each of the two test sets.
- describe in your own words what the images look like in the two test sets.
- Include plots of the four basis elements of each of the three bases and describe shortly the visual differences between the three bases.
- Print the mean of the error norms for the six combinations in a table (two test sets against the three bases).
- Which basis works best for test set 1? Why? Explain your reasoning!
- Which basis works best for test set 2? Why? Explain your reasoning!